

SIT330-770: Natural Language Processing

Week 6 - Vector Semantics and Embeddings

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of
Sci Eng & Built Env

reda.bouadjenek@deakin.edu.au

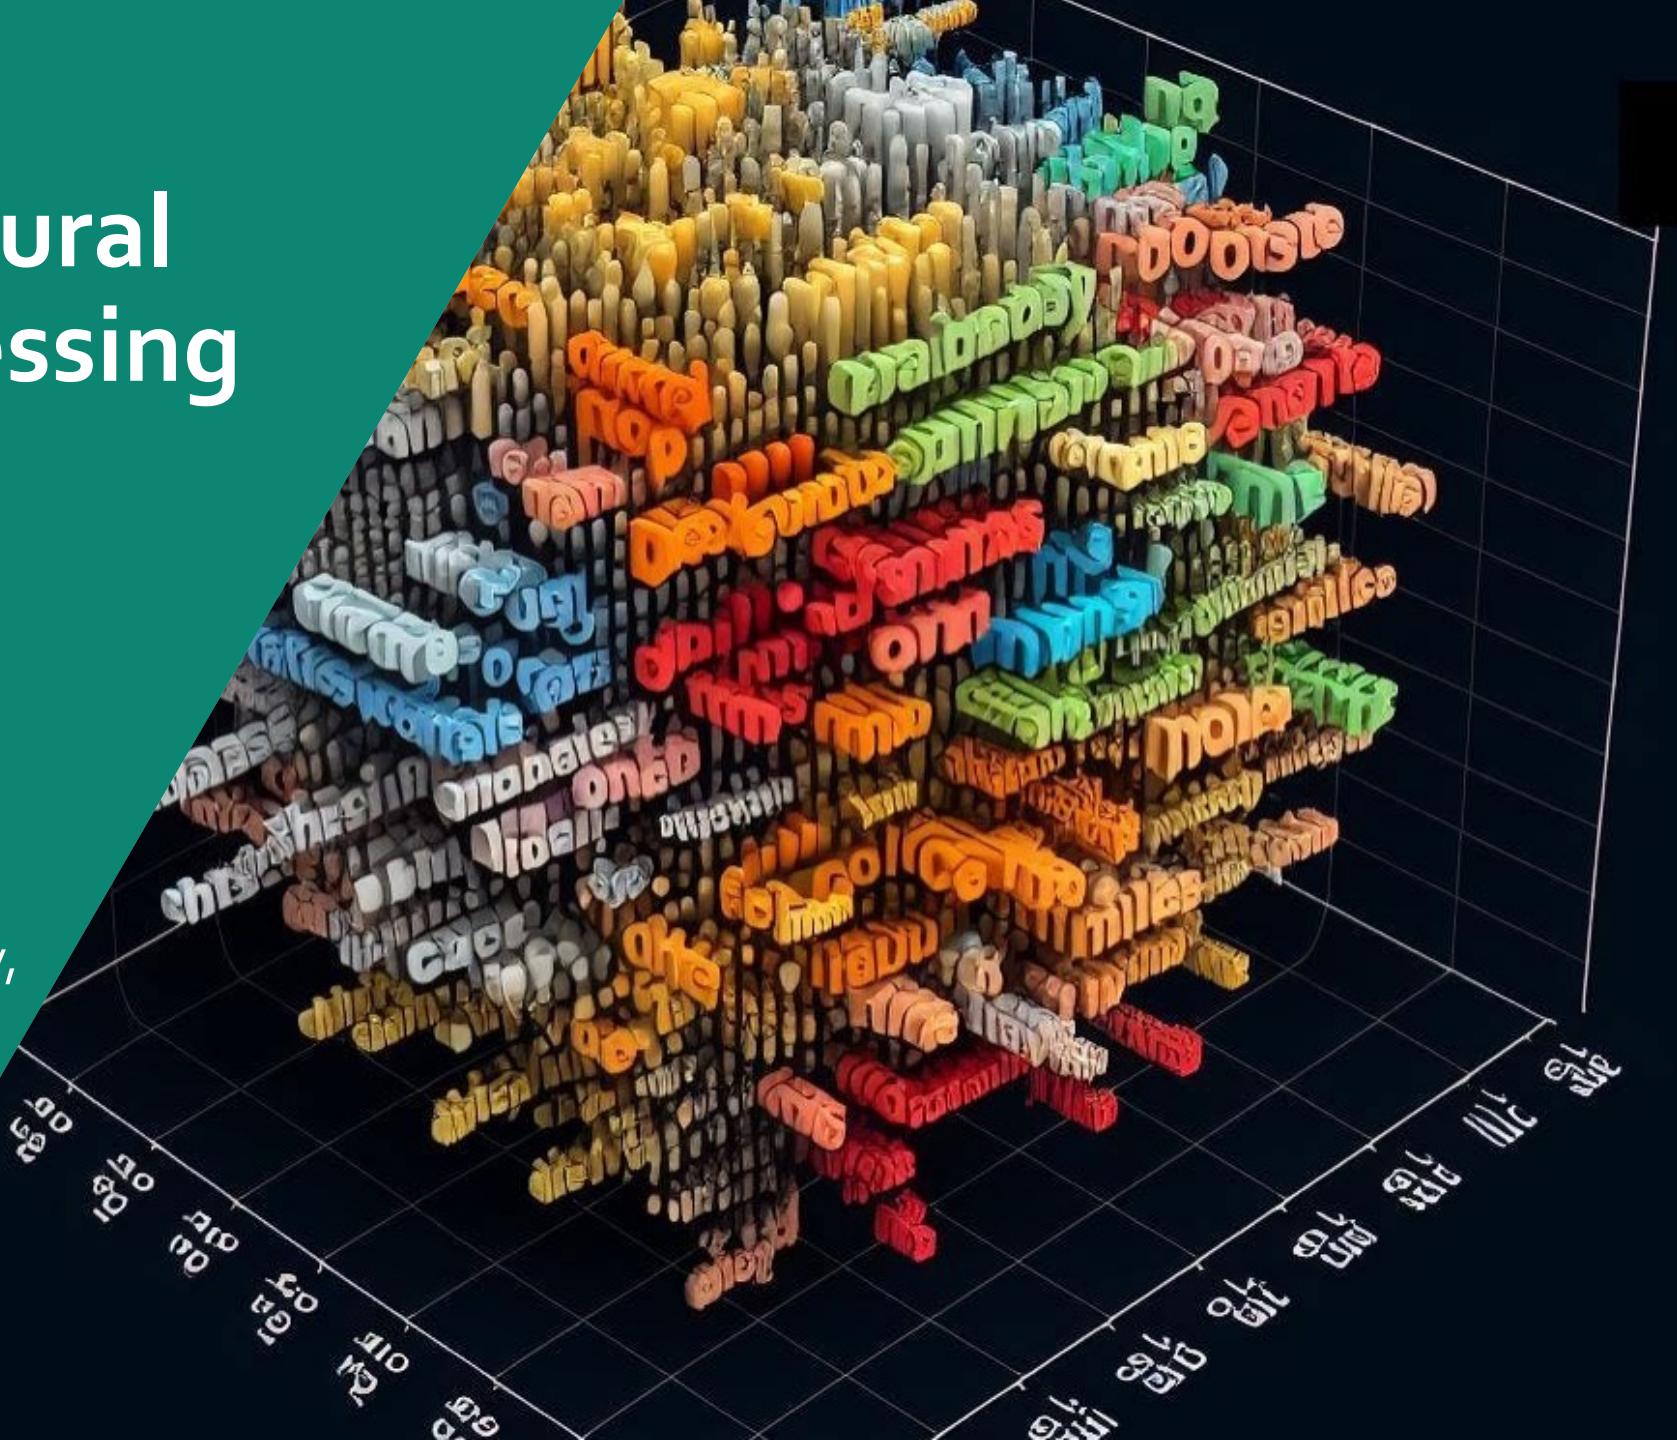


SIT330-770: Natural Language Processing

Week 6.1 - Word Meaning

Dr. Mohamed Reda Bouadjenek

School of Information Technology,
Faculty of Sci Eng & Built Env

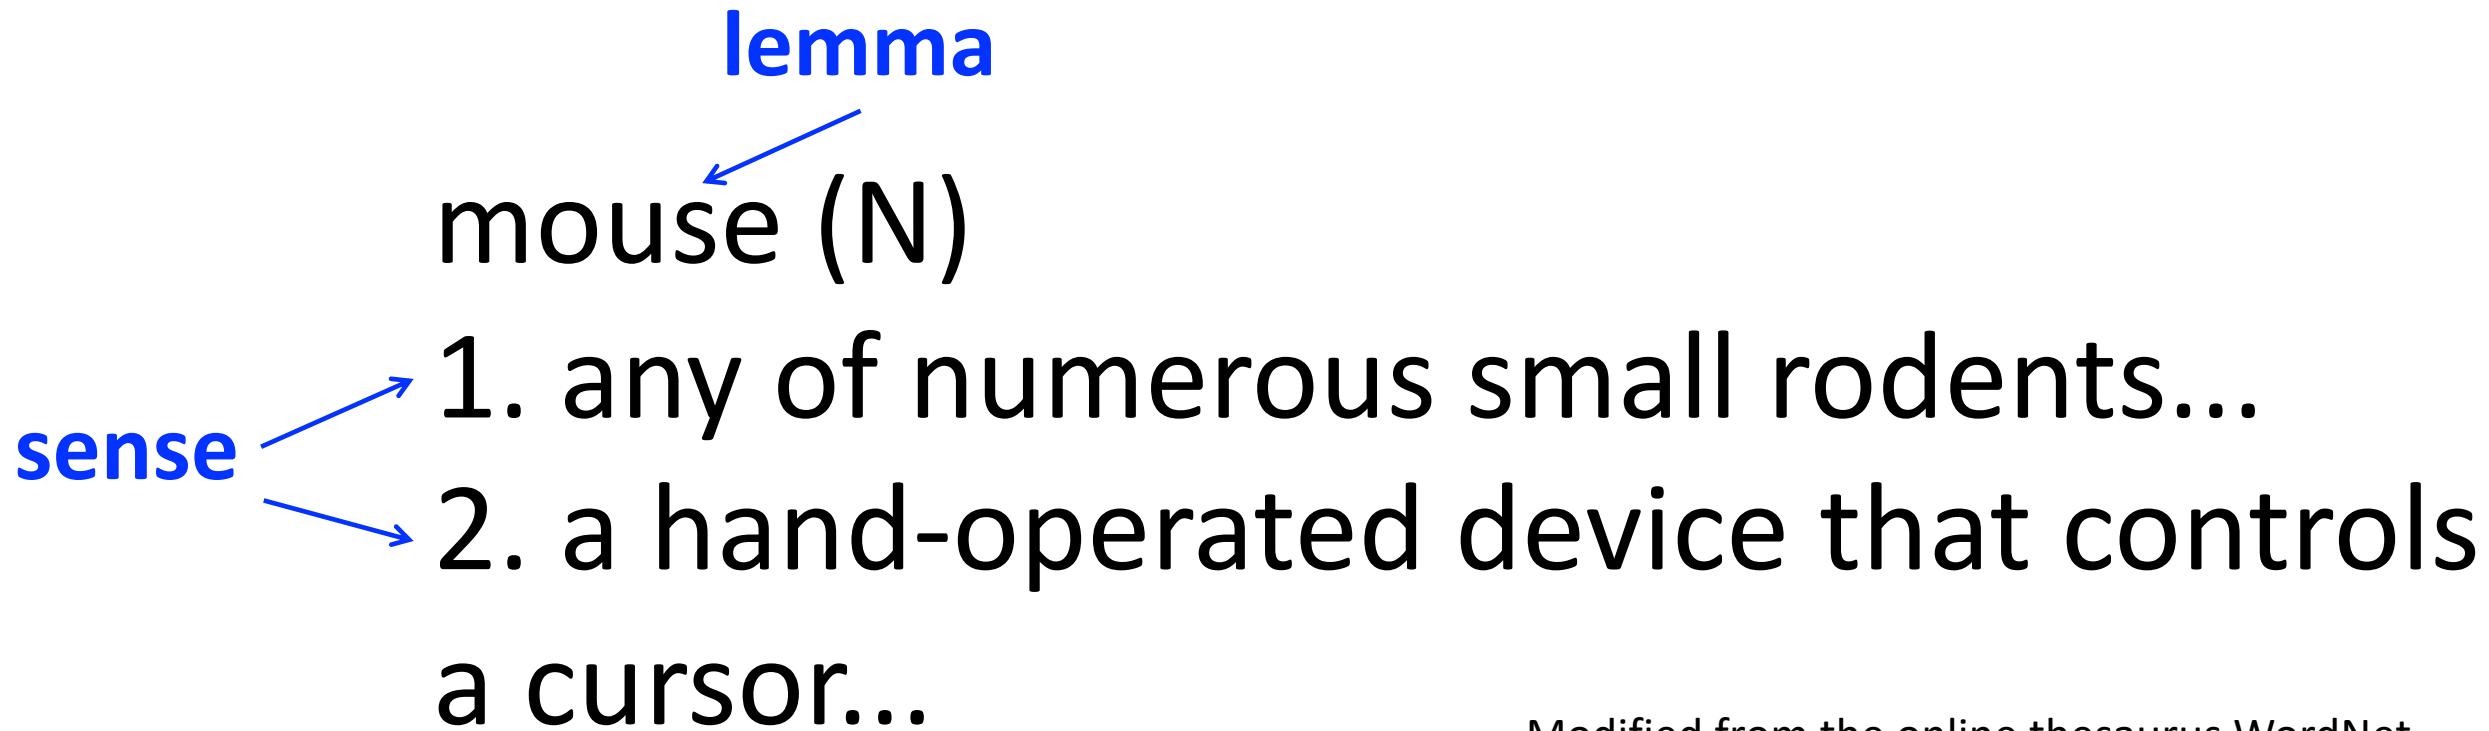


What do words mean?



- N-gram or text classification methods we've seen so far
 - Words are just strings (or indices w_i in a vocabulary list)
 - That's not very satisfactory!
- Introductory logic classes:
 - The meaning of "dog" is DOG; cat is CAT
$$\forall x \text{DOG}(x) \rightarrow \text{MAMMAL}(x)$$
- Old linguistics joke by Barbara Partee in 1967:
 - Q: What's the meaning of life?
 - A: LIFE
- That seems hardly better!

- What should a theory of word meaning do for us?
- Let's look at some desiderata
- From **lexical semantics**, the linguistic study of word meaning



Modified from the online thesaurus WordNet

A **sense** or “**concept**” is the meaning component of a word
Lemmas can be **polysemous** (have multiple senses)

Relations between senses: Synonymy



- Synonyms have the same meaning in some or all contexts.
 - filbert / hazelnut
 - couch / sofa
 - big / large
 - automobile / car
 - vomit / throw up
 - water / H₂O

- Note that there are probably no examples of perfect synonymy.
 - Even if many aspects of meaning are identical
 - Still may differ based on politeness, slang, register, genre, etc.

Relation: Synonymy?



water/ H_2O

" H_2O " in a surfing guide?

big/large

my big sister != my large sister

The Linguistic Principle of Contrast



- Difference in form → difference in meaning

Re: "exact" synonyms

je ne crois pas qu'il y ait de mot synonyme dans aucune Langue. Je le dis par con-

[I do not believe that there is a synonymous word in any language]

Thanks to Mark Aronoff!

LA JUSTESSE
DE LA
LANGUE FRANÇOISE,
ou
LES DIFFÉRENTES SIGNIFICATIONS
DES MOTS QUI PASSENT
POUR
SYNONIMES.

Par M. l'Abbé GIRARD C. D. M. D. D. B.



A PARIS,
Chez LAURENT D'HOURY, Imprimeur-
Libraire, au bas de la rue de la Harpe, vis-
à vis la rue S. Severin, au Saint-Esprit.

M. DCC. XVIII.
Avec Approbation & Privilegs du Roy.

Words with similar meanings. Not synonyms, but sharing some element of meaning

car, bicycle

cow, horse

Ask humans how similar 2 words are



word1	word2	similarity
vanish	disappear	9.8
behave	obey	7.3
belief	impression	5.95
muscle	bone	3.65
modest	flexible	0.98
hole	agreement	0.3

- Also called "word association"
- Words can be related in any way, perhaps via a semantic frame or field
 - coffee, tea: **similar**
 - coffee, cup: **related**, not similar

- Words that
 - cover a particular semantic domain
 - bear structured relations with each other.

hospitals

surgeon, scalpel, nurse, anaesthetic, hospital

restaurants

waiter, menu, plate, food, menu, chef

houses

door, roof, kitchen, family, bed

- Senses that are opposites with respect to only one feature of meaning
- Otherwise, they are very similar!

dark/light short/long fast/slow rise/fall

hot/cold up/down in/out

- More formally: antonyms can
 - define a binary opposition or be at opposite ends of a scale
 - long/short, fast/slow
 - Be *reversives*:
 - rise/fall, up/down

- Words have **affective** meanings
 - Positive connotations (*happy*)
 - Negative connotations (*sad*)
- Connotations can be subtle:
 - Positive connotation: *copy, replica, reproduction*
 - Negative connotation: *fake, knockoff, forgery*
- Evaluation (sentiment!)
 - Positive evaluation (*great, love*)
 - Negative evaluation (*terrible, hate*)

Osgood et al. (1957)

- Words seem to vary along 3 affective dimensions:
 - **valence**: the pleasantness of the stimulus
 - **arousal**: the intensity of emotion provoked by the stimulus
 - **dominance**: the degree of control exerted by the stimulus

	Word	Score		Word	Score
Valence	love	1.000		toxic	0.008
	happy	1.000		nightmare	0.005
Arousal	elated	0.960		mellow	0.069
	frenzy	0.965		napping	0.046
Dominance	powerful	0.991		weak	0.045
	leadership	0.983		empty	0.081

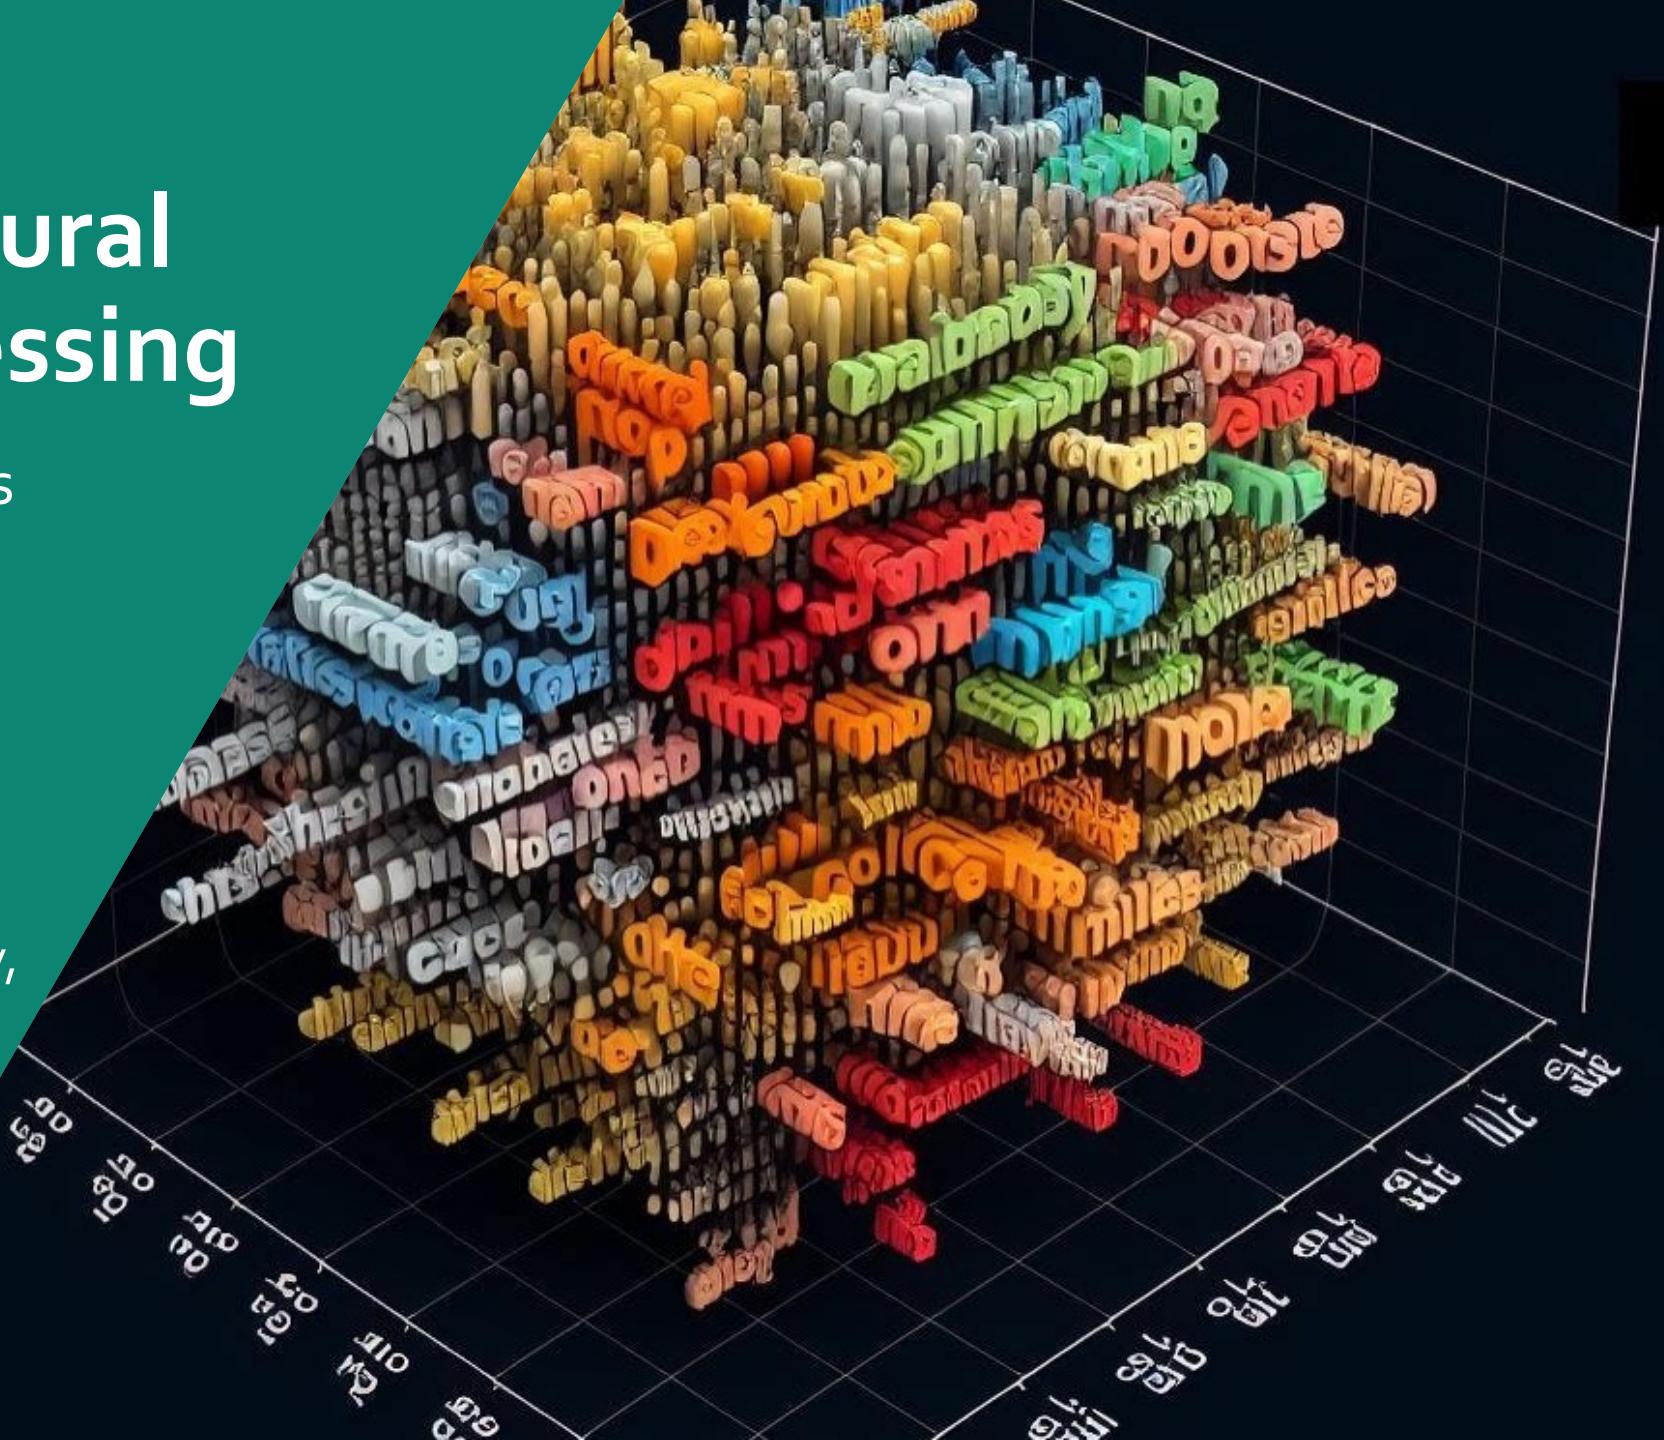
- **Concepts or word senses**
 - Have a complex many-to-many association with **words** (homonymy, multiple senses)
- Have relations with each other
 - Synonymy
 - Antonymy
 - Similarity
 - Relatedness
 - Connotation

SIT330-770: Natural Language Processing

Week 6.2 - Vector Semantics

Dr. Mohamed Reda Bouadjenek

School of Information Technology,
Faculty of Sci Eng & Built Env



- Can we build a theory of how to represent word meaning, that accounts for at least some of the desiderata?
- We'll introduce **vector semantics**
 - The standard model in language processing!
 - Handles many of our goals!

- PI #43:

"The meaning of a word is its use in the language"

Let's define words by their usages



- One way to define "usage":
- words are defined by their environments (the words around them)
- Zellig Harris (1954):
- **If A and B have almost identical environments we say that they are synonyms.**

What does recent English borrowing *ongchoi* mean?



- Suppose you see these sentences:
 - Ong choi is delicious **sautéed with garlic**.
 - Ong choi is superb **over rice**
 - Ong choi **leaves** with salty sauces
- And you've also seen these:
 - ...spinach **sautéed with garlic over rice**
 - Chard stems and **leaves** are **delicious**
 - Collard greens and other **salty** leafy greens
- Conclusion:
 - Ongchoi is a leafy green like spinach, chard, or collard greens
 - We could conclude this based on words like "leaves" and "delicious" and "sautéed"

空心菜
kangkong
rau muống
...



Yamaguchi, Wikimedia Commons, public domain

Idea 1: Defining meaning by linguistic distribution



- Let's define the meaning of a word by its distribution in language use, meaning its neighboring words or grammatical environments.

Idea 2: Meaning as a point in space (Osgood et al. 1957)



- 3 affective dimensions for a word
 - **valence**: pleasantness
 - **arousal**: intensity of emotion
 - **dominance**: the degree of control exerted

	Word	Score		Word	Score
Valence	love	1.000		toxic	0.008
	happy	1.000		nightmare	0.005
Arousal	elated	0.960		mellow	0.069
	frenzy	0.965		napping	0.046
Dominance	powerful	0.991		weak	0.045
	leadership	0.983		empty	0.081

NRC VAD Lexicon
(Mohammad 2018)

- Hence the connotation of a word is a vector in 3-space

Idea 1: Defining meaning by linguistic distribution

Idea 2: Meaning as a point in multidimensional space

Defining meaning as a point in space based on distribution



- Each word = a vector (not just "good" or " w_{45} ")
- Similar words are "**nearby in semantic space**"
- We build this space automatically by seeing which words are **nearby in text**



- Called an "embedding" because it's embedded into a space (see textbook)
- The standard way to represent meaning in NLP
- **Every modern NLP algorithm uses embeddings as the representation of word meaning**
- Fine-grained model of meaning for similarity

- Consider sentiment analysis:
 - With **words**, a feature is a word identity
 - Feature 5: 'The previous word was "terrible"
 - requires **exact same word** to be in training and test
 - With **embeddings**:
 - Feature is a word vector
 - 'The previous word was vector [35, 22, 17...]
 - Now in the test set we might see a similar vector [34, 21, 14]
 - We can generalize to **similar but unseen** words!!!

We'll discuss 2 kinds of embeddings



- **tf-idf**
 - Information Retrieval workhorse!
 - A common baseline model
 - **Sparse** vectors
 - Words are represented by (a simple function of) the **counts** of nearby words
- **Word2vec**
 - **Dense** vectors
 - Representation is created by training a classifier to **predict** whether a word is likely to appear nearby
 - Later we'll discuss extensions called **contextual embeddings**

From now on: Computing with meaning representations instead of string representations



荃者所以在鱼，得鱼而忘荃

Nets are for fish;

言者所以在意，得意而忘言

Once you get the fish, you can forget the net.

Words are for meaning;

Once you get the meaning, you can forget the words

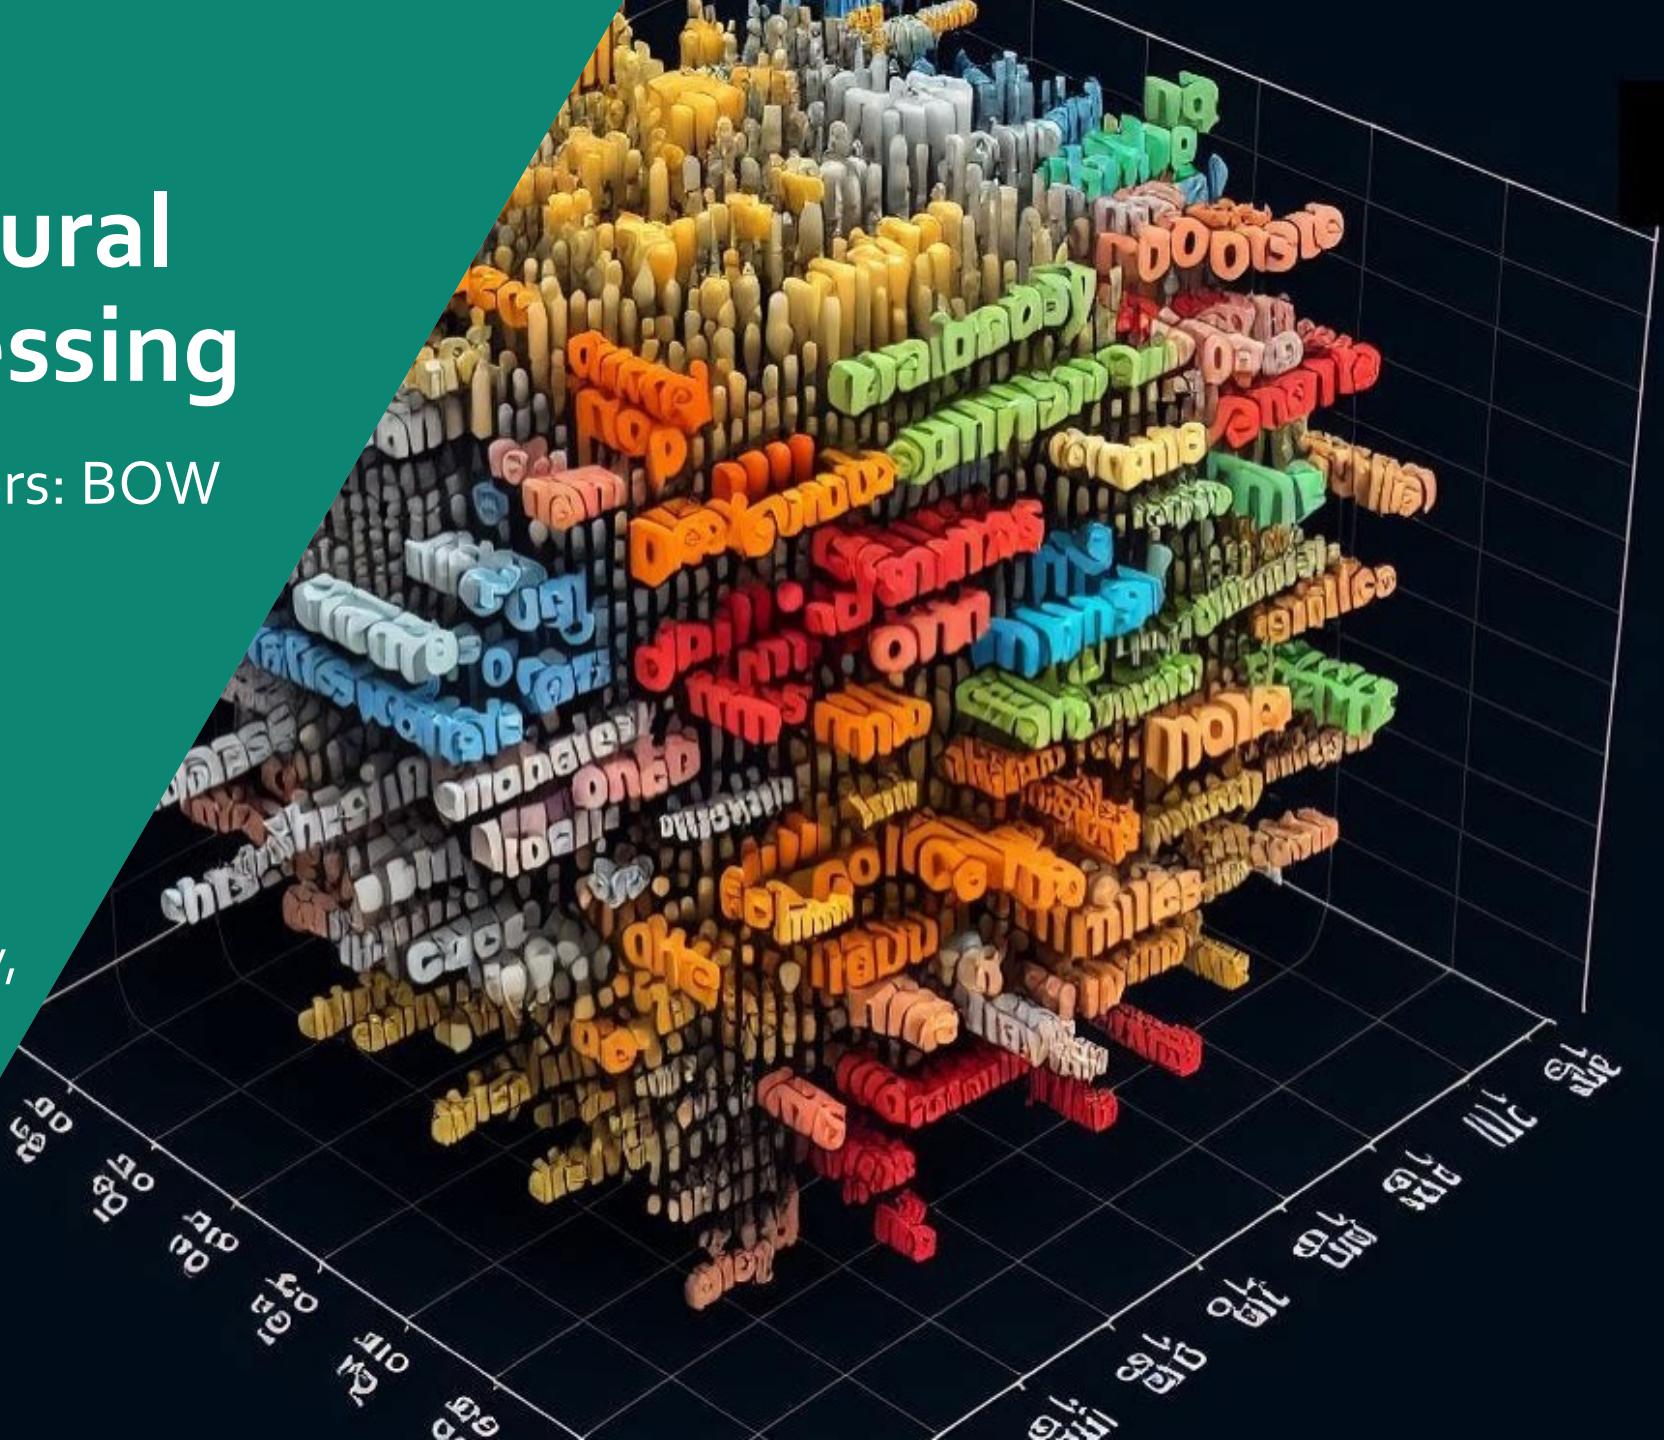
庄子(Zhuangzi), Chapter 26

SIT330-770: Natural Language Processing

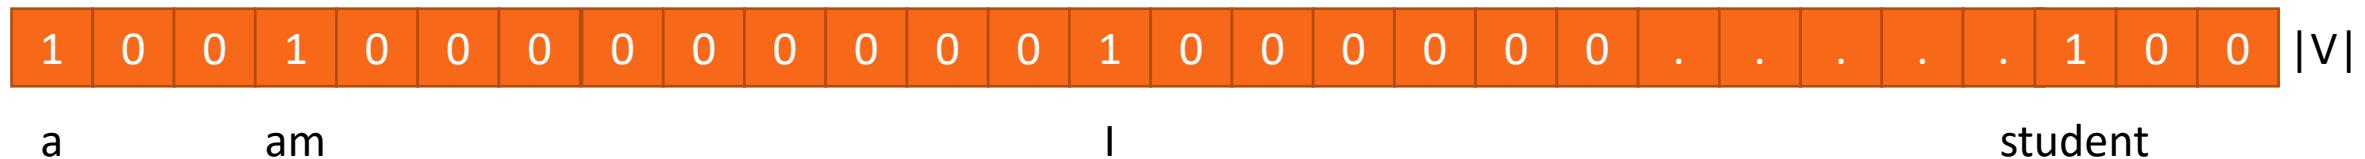
Week 6.3 – Words and Vectors: BOW

Dr. Mohamed Reda Bouadjenek

School of Information Technology,
Faculty of Sci Eng & Built Env



- A document is represented as vector of words.
 - One dimension per word.
 - Vector size is the vocabulary size, e.g., English may contain 100k words.
 - Different weighting schemas can be used, e.g., tf, log(tf), tf-idf, Boolean, etc.
 - Sparse vector, e.g., almost all values are zeros.



Order matters for NLP tasks!



- Assumes independence between words:
 - The sentences “**John likes Mary**” has the same representation as “**Mary likes John**” – even though the semantic is different).
- May work well for Information Retrieval tasks, but not for NLP tasks!
 - Sentiment analysis:
“Ah **no**, there are good movies on Netflix!” vs. “Ah, there are **no** good movies on Netflix!”

Order matters for NLP tasks!



“Ah **no**, there are good movies on Netflix!” vs. “Ah, there are **no** good movie on Netflix!”

- Use N-grams.
 - Dimensionality grows exponentially V^N .
 - 3-grams with English: $(10^5)^3 = 10^{15} = 1,000,000,000,000,000$ entries.
 - Too expensive!

- The dot product between two vectors is a scalar:

$$\text{dot product}(\mathbf{v}, \mathbf{w}) = \mathbf{v} \cdot \mathbf{w} = \sum_{i=1}^N v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

- The dot product tends to be high when the two vectors have large values in the same dimensions
- Dot product can thus be a useful similarity metric between vectors

Problem with raw dot-product



- Dot product favors long vectors
 - Dot product is higher if a vector is longer (has higher values in many dimension)
 - Vector length:
-
- Frequent words (of, the, you) have long vectors (since they occur many times with other words).
 - So dot product overly favors frequent words

Alternative: cosine for computing word similarity



$$\text{cosine}(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}| |\vec{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

Based on the definition of the dot product between two vectors a and b

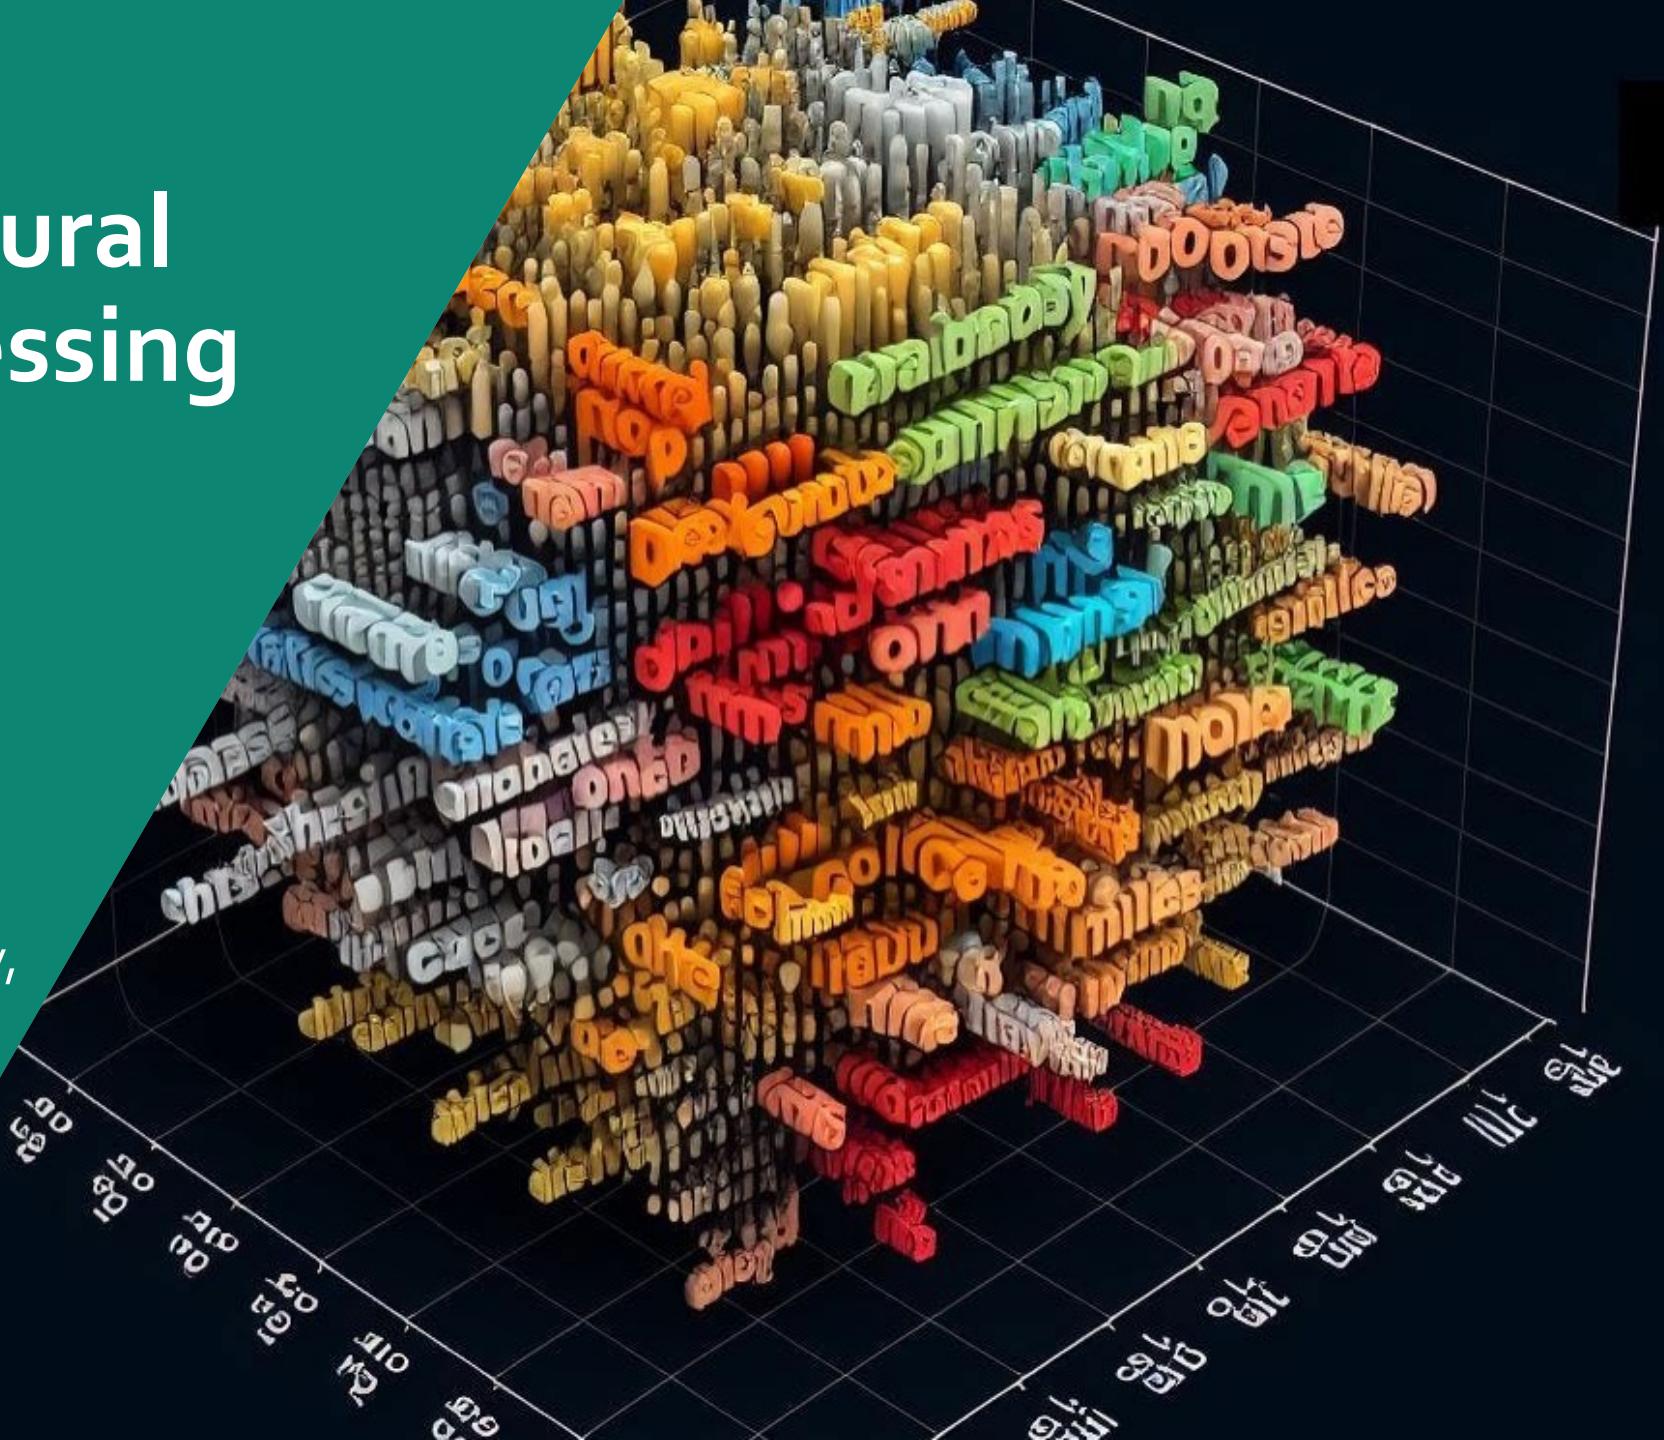
$$\begin{aligned}\mathbf{a} \cdot \mathbf{b} &= |\mathbf{a}| |\mathbf{b}| \cos \theta \\ \frac{\mathbf{a} \cdot \mathbf{b}}{|\mathbf{a}| |\mathbf{b}|} &= \cos \theta\end{aligned}$$

SIT330-770: Natural Language Processing

Week 6.4 – Word2Vec

Dr. Mohamed Reda Bouadjenek

School of Information Technology,
Faculty of Sci Eng & Built Env



- tf-idf (or PMI) vectors are
 - **long** (length $|V| = 20,000$ to $50,000$)
 - **sparse** (most elements are zero)
- Alternative: learn vectors which are
 - **short** (length 50-1000)
 - **dense** (most elements are non-zero)

- Why dense vectors?
 - Short vectors may be easier to use as **features** in machine learning (fewer weights to tune)
 - Dense vectors may **generalize** better than explicit counts
 - Dense vectors may do better at capturing synonymy:
 - *car* and *automobile* are synonyms; but are distinct dimensions
 - a word with *car* as a neighbor and a word with *automobile* as a neighbor should be similar, but aren't
 - **In practice, they work better**

Common methods for getting short dense vectors



- “Neural Language Model”-inspired models
 - Word2vec (skipgram, CBOW), GloVe
- Singular Value Decomposition (SVD)
 - A special case of this is called LSA – Latent Semantic Analysis
- Alternative to these "static embeddings":
 - Contextual Embeddings (ELMo, BERT)
 - Compute distinct embeddings for a word in its context
 - Separate embeddings for each token of a word

Simple static embeddings you can download!



- Word2vec (Mikolov et al)
- <https://code.google.com/archive/p/word2vec/>
- GloVe (Pennington, Socher, Manning)
- <http://nlp.stanford.edu/projects/glove/>

- Popular embedding method
- Very fast to train
- Code available on the web
- Idea: **predict** rather than **count**
- Word2vec provides various options. We'll do:
 - **skip-gram with negative sampling (SGNS)**

- Instead of **counting** how often each word w occurs near "apricot"
 - Train a classifier on a binary **prediction** task:
 - Is w likely to show up near "apricot"?
- We don't actually care about this task
 - But we'll take the learned classifier weights as the word embeddings
- Big idea: **self-supervision**:
 - A word c that occurs near apricot in the corpus acts as the gold "correct answer" for supervised learning
 - No need for human labels
 - Bengio et al. (2003); Collobert et al. (2011)

Approach: predict if candidate word c is a "neighbor"



1. Treat the target word t and a neighboring context word c as **positive examples**.
2. Randomly sample other words in the lexicon to get negative examples
3. Use logistic regression to train a classifier to distinguish those two cases
4. Use the learned weights as the embeddings

- (assuming a +/- 2 word window)

...lemon, a [tablespoon of apricot jam, a] pinch...

c1 c2 [target] c3 c4

- Goal: train a classifier that is given a candidate (**word, context**) pair

(apricot, jam)

(apricot, aardvark)

...

- And assigns each pair a probability:

- $P(+|w, c)$
 - $P(-|w, c) = 1 - P(+|w, c)$

Similarity is computed from dot product



- Remember: two vectors are similar if they have a high dot product
 - Cosine is just a normalized dot product
- So:
 - $\text{Similarity}(w, c) \propto w \cdot c$
- We'll need to normalize to get a probability
 - (cosine isn't a probability either)

- $\text{Sim}(w, c) \approx w \cdot c$
- To turn this into a probability
- We'll use the sigmoid from logistic regression:

$$P(+|w, c) = \sigma(c \cdot w) = \frac{1}{1 + \exp(-c \cdot w)}$$

$$\begin{aligned} P(-|w, c) &= 1 - P(+|w, c) \\ &= \sigma(-c \cdot w) = \frac{1}{1 + \exp(c \cdot w)} \end{aligned}$$

How Skip-Gram Classifier computes $P(+|w, c)$



$$P(+|w, c) = \sigma(c \cdot w) = \frac{1}{1 + \exp(-c \cdot w)}$$

- This is for one context word, but we have lots of context words.
- We'll assume independence and just multiply them:

$$P(+|w, c_{1:L}) = \prod_{i=1}^L \sigma(c_i \cdot w)$$

$$\log P(+|w, c_{1:L}) = \sum_{i=1}^L \log \sigma(c_i \cdot w)$$

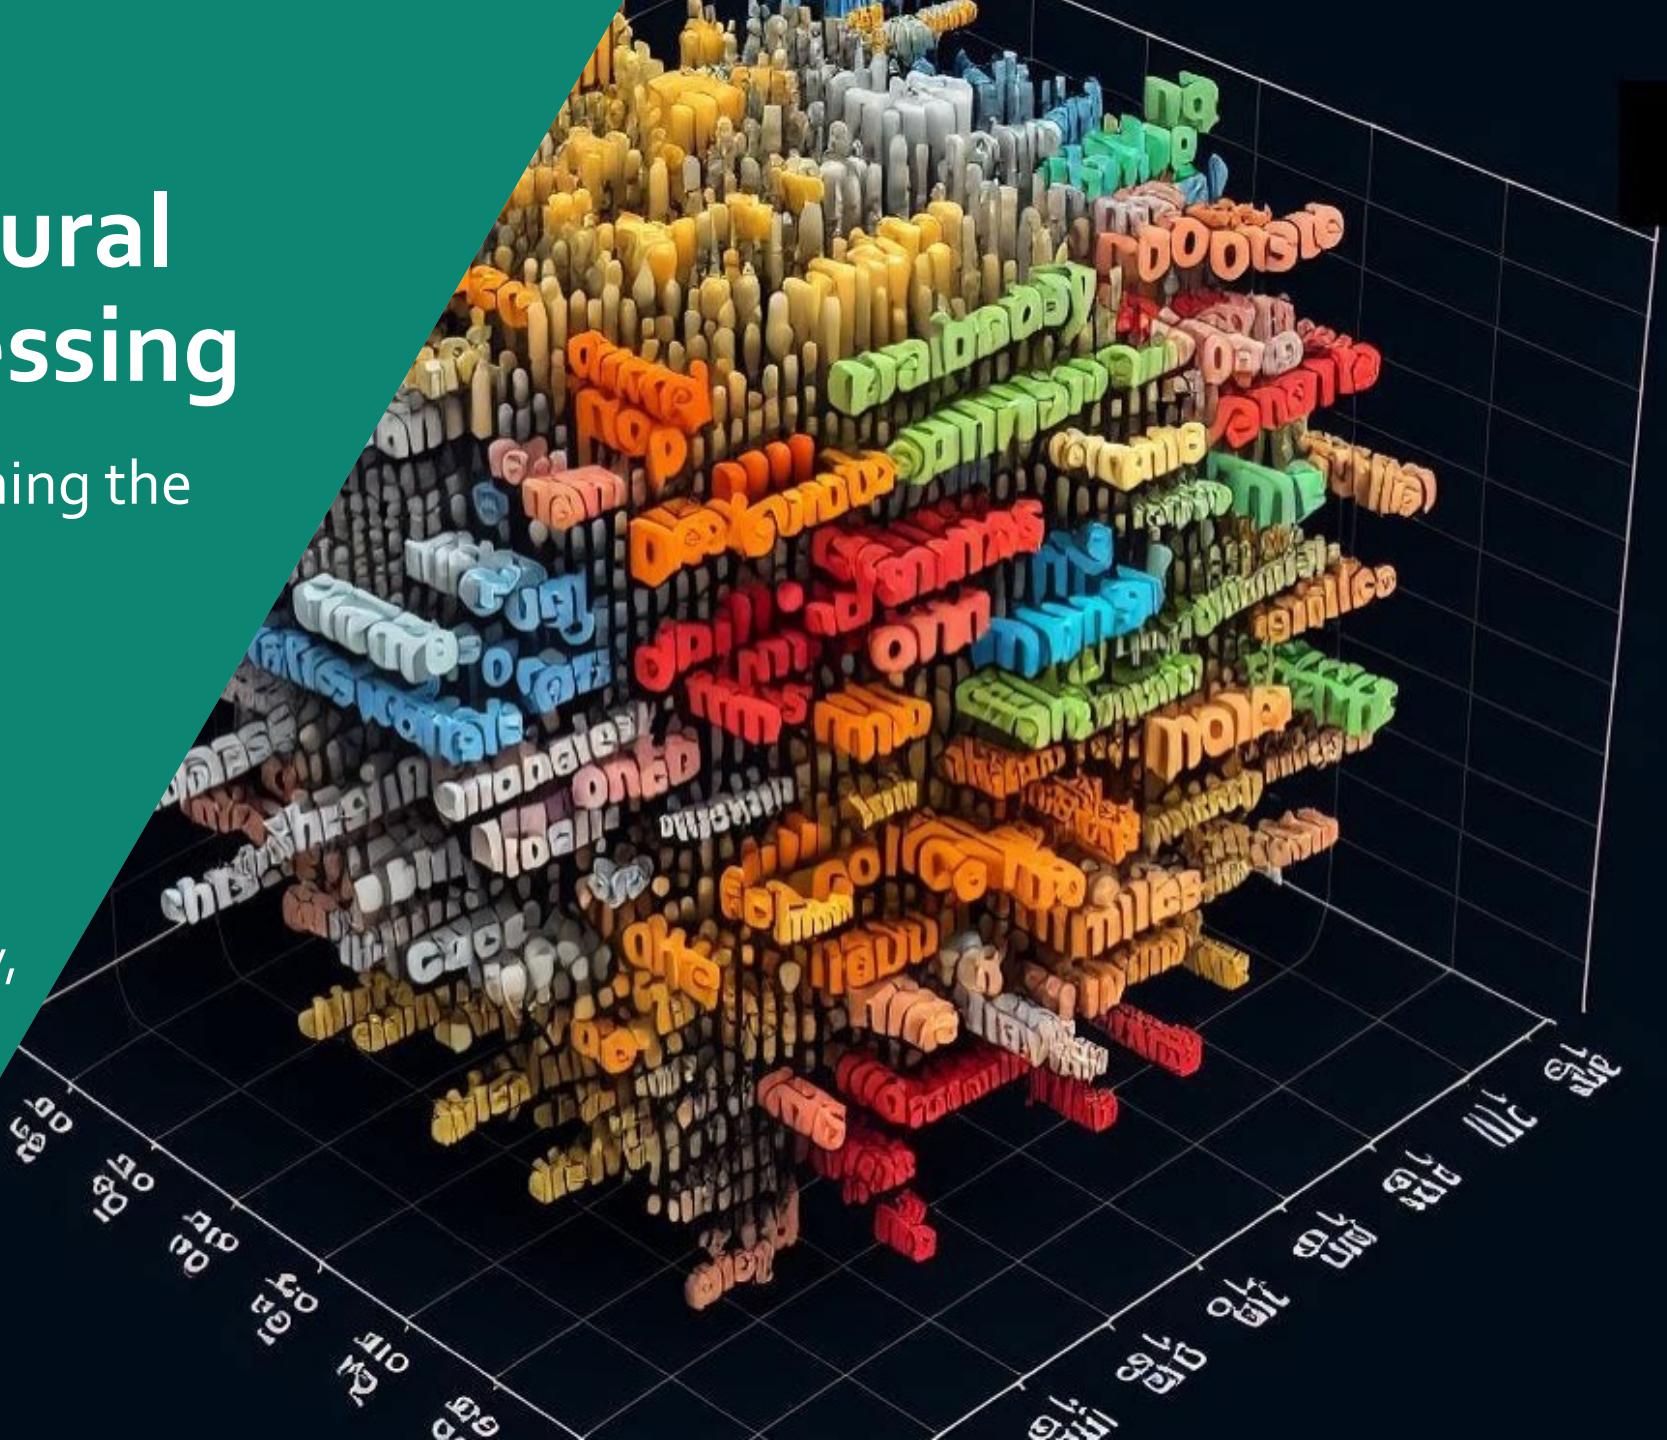
- A probabilistic classifier, given
 - a test target word w
 - its context window of L words $c_{1:L}$
- Estimates probability that w occurs in this window based on similarity of w (embeddings) to $C_{1:L}$ (embeddings).
- To compute this, we just need embeddings for all the words.

SIT330-770: Natural Language Processing

Week 6.5 – Word2vec: Learning the embeddings

Dr. Mohamed Reda Bouadjenek

School of Information Technology,
Faculty of Sci Eng & Built Env



...lemon, a [tablespoon of apricot jam, a] pinch...

c1

c2 [target]

c3

c4



positive examples +

t c

apricot tablespoon

apricot of

apricot jam

apricot a

...lemon, a [tablespoon of apricot jam, a] pinch...

c1

c2 [target]

c3

c4



positive examples +

t c

apricot tablespoon

apricot of

apricot jam

apricot a

For each positive example we'll grab k negative examples, sampling by frequency

...lemon, a [tablespoon of apricot jam, a] pinch...

c1	t	c
	apricot	tablespoon
	apricot	of
	apricot	jam
	apricot	a

positive examples +

negative examples -

c2 [target]	t	c3	c4
	apricot	aardvark	apricot
	apricot	my	seven
	apricot	where	forever
	apricot	coaxial	dear
			if

- Given the set of positive and negative training instances, and an initial set of embedding vectors
- The goal of learning is to adjust those word vectors such that we:
 - **Maximize** the similarity of the target word, context word pairs (w, c_{pos}) drawn from the positive data
 - **Minimize** the similarity of the (w, c_{neg}) pairs drawn from the negative data.

Loss function for one w with $c_{pos}, c_{neg_1} \dots c_{neg_k}$

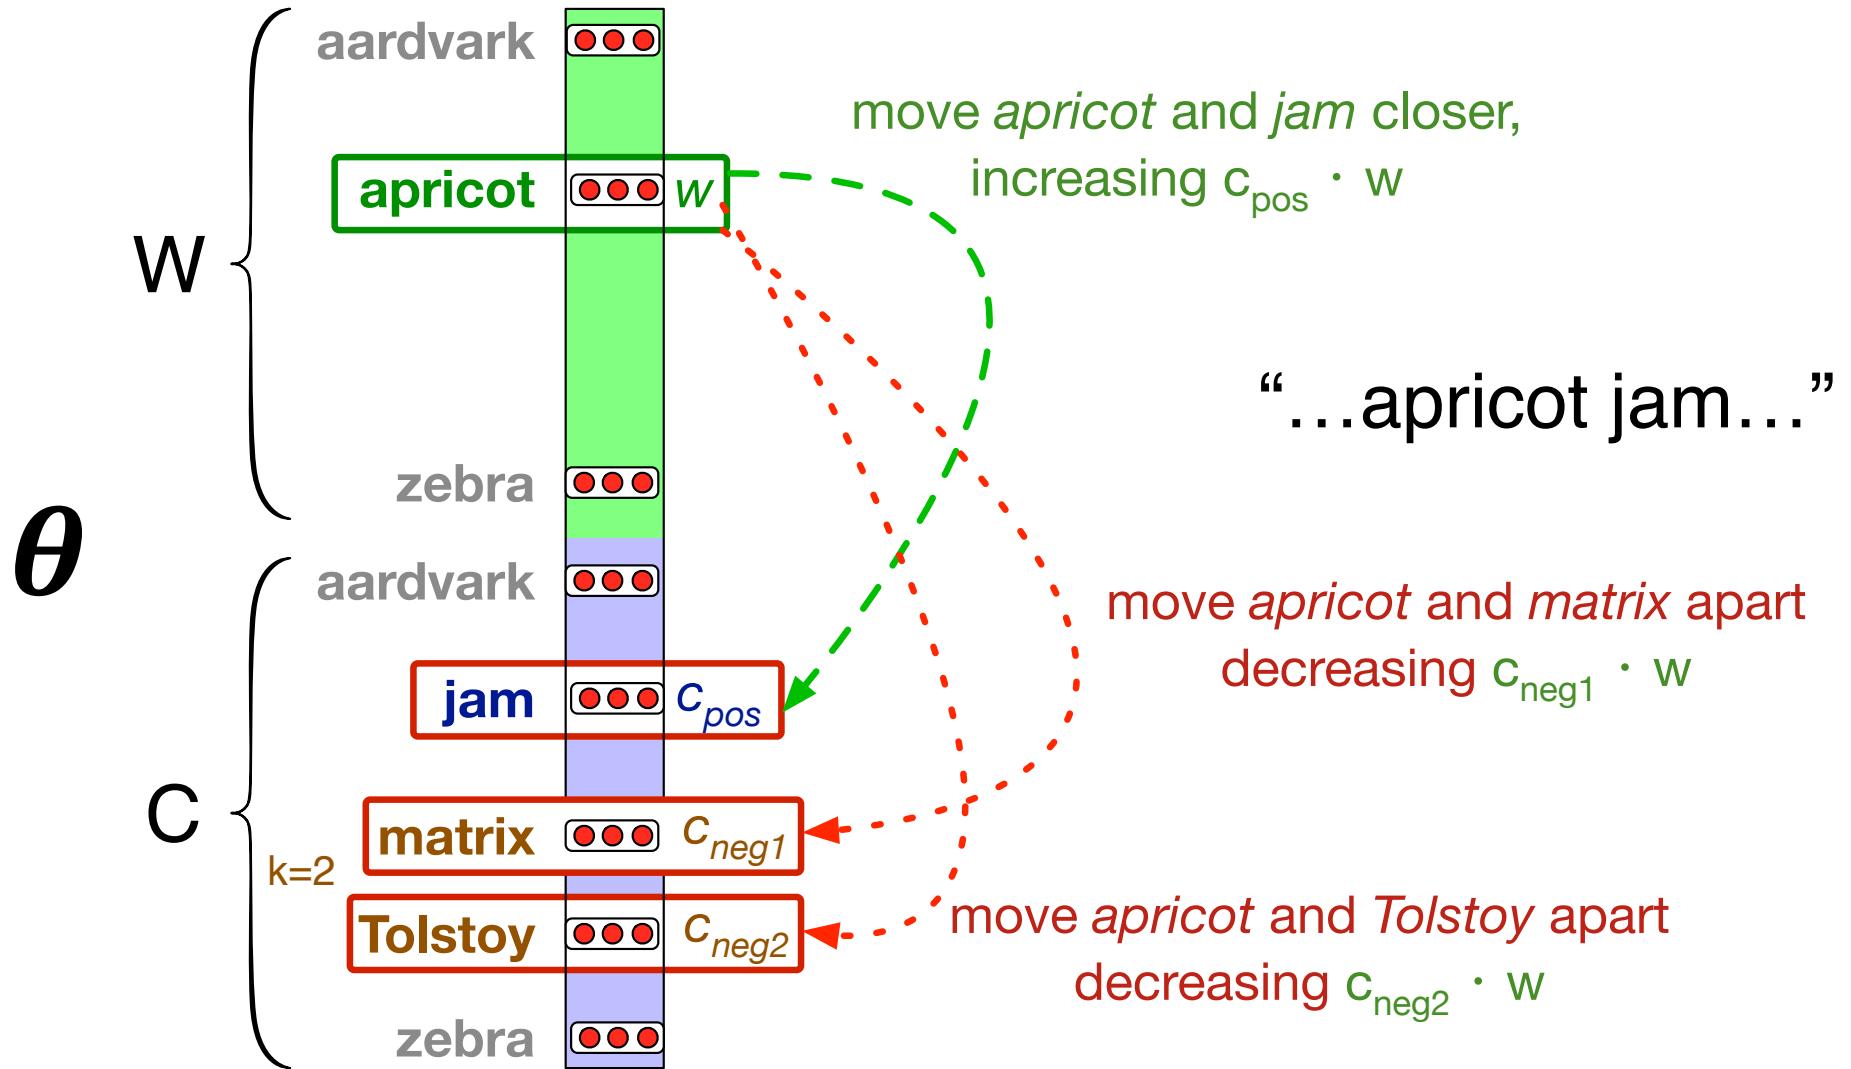


- Maximize the similarity of the target with the actual context words, and minimize the similarity of the target with the k negative sampled non-neighbor words.

$$\begin{aligned} L_{CE} &= -\log \left[P(+|w, c_{pos}) \prod_{i=1}^k P(-|w, c_{neg_i}) \right] \\ &= - \left[\log P(+|w, c_{pos}) + \sum_{i=1}^k \log P(-|w, c_{neg_i}) \right] \\ &= - \left[\log P(+|w, c_{pos}) + \sum_{i=1}^k \log (1 - P(+|w, c_{neg_i})) \right] \\ &= - \left[\log \sigma(c_{pos} \cdot w) + \sum_{i=1}^k \log \sigma(-c_{neg_i} \cdot w) \right] \end{aligned}$$

- How to learn?
 - Stochastic gradient descent!
- We'll adjust the word weights to
 - make the positive pairs more likely
 - and the negative pairs less likely,
 - over the entire training set.

Intuition of one step of gradient descent



- At each step
 - Direction: We move in the reverse direction from the gradient of the loss function
 - Magnitude: we move the value of this gradient $\frac{d}{dw} L(f(x; w), y)$ weighted by a **learning rate** η
 - Higher learning rate means move w faster

$$w^{t+1} = w^t - \eta \frac{d}{dw} L(f(x; w), y)$$

The derivatives of the loss function



$$L_{CE} = - \left[\log \sigma(c_{pos} \cdot w) + \sum_{i=1}^k \log \sigma(-c_{neg_i} \cdot w) \right]$$

$$\frac{\partial L_{CE}}{\partial c_{pos}} = [\sigma(c_{pos} \cdot w) - 1]w$$

$$\frac{\partial L_{CE}}{\partial c_{neg}} = [\sigma(c_{neg} \cdot w)]w$$

$$\frac{\partial L_{CE}}{\partial w} = [\sigma(c_{pos} \cdot w) - 1]c_{pos} + \sum_{i=1}^k [\sigma(c_{neg_i} \cdot w)]c_{neg_i}$$

- Start with randomly initialized C and W matrices, then incrementally do updates

$$c_{pos}^{t+1} = c_{pos}^t - \eta [\sigma(c_{pos}^t \cdot w^t) - 1]w^t$$

$$c_{neg}^{t+1} = c_{neg}^t - \eta [\sigma(c_{neg}^t \cdot w^t)]w^t$$

$$w^{t+1} = w^t - \eta \left[[\sigma(c_{pos} \cdot w^t) - 1]c_{pos} + \sum_{i=1}^k [\sigma(c_{neg_i} \cdot w^t)]c_{neg_i} \right]$$

Two sets of embeddings



- SGNS learns two sets of embeddings
 - Target embeddings matrix W
 - Context embedding matrix C
- It's common to just add them together, representing word i as the vector $w_i + c_i$

Summary: How to learn word2vec (skip-gram) embeddings



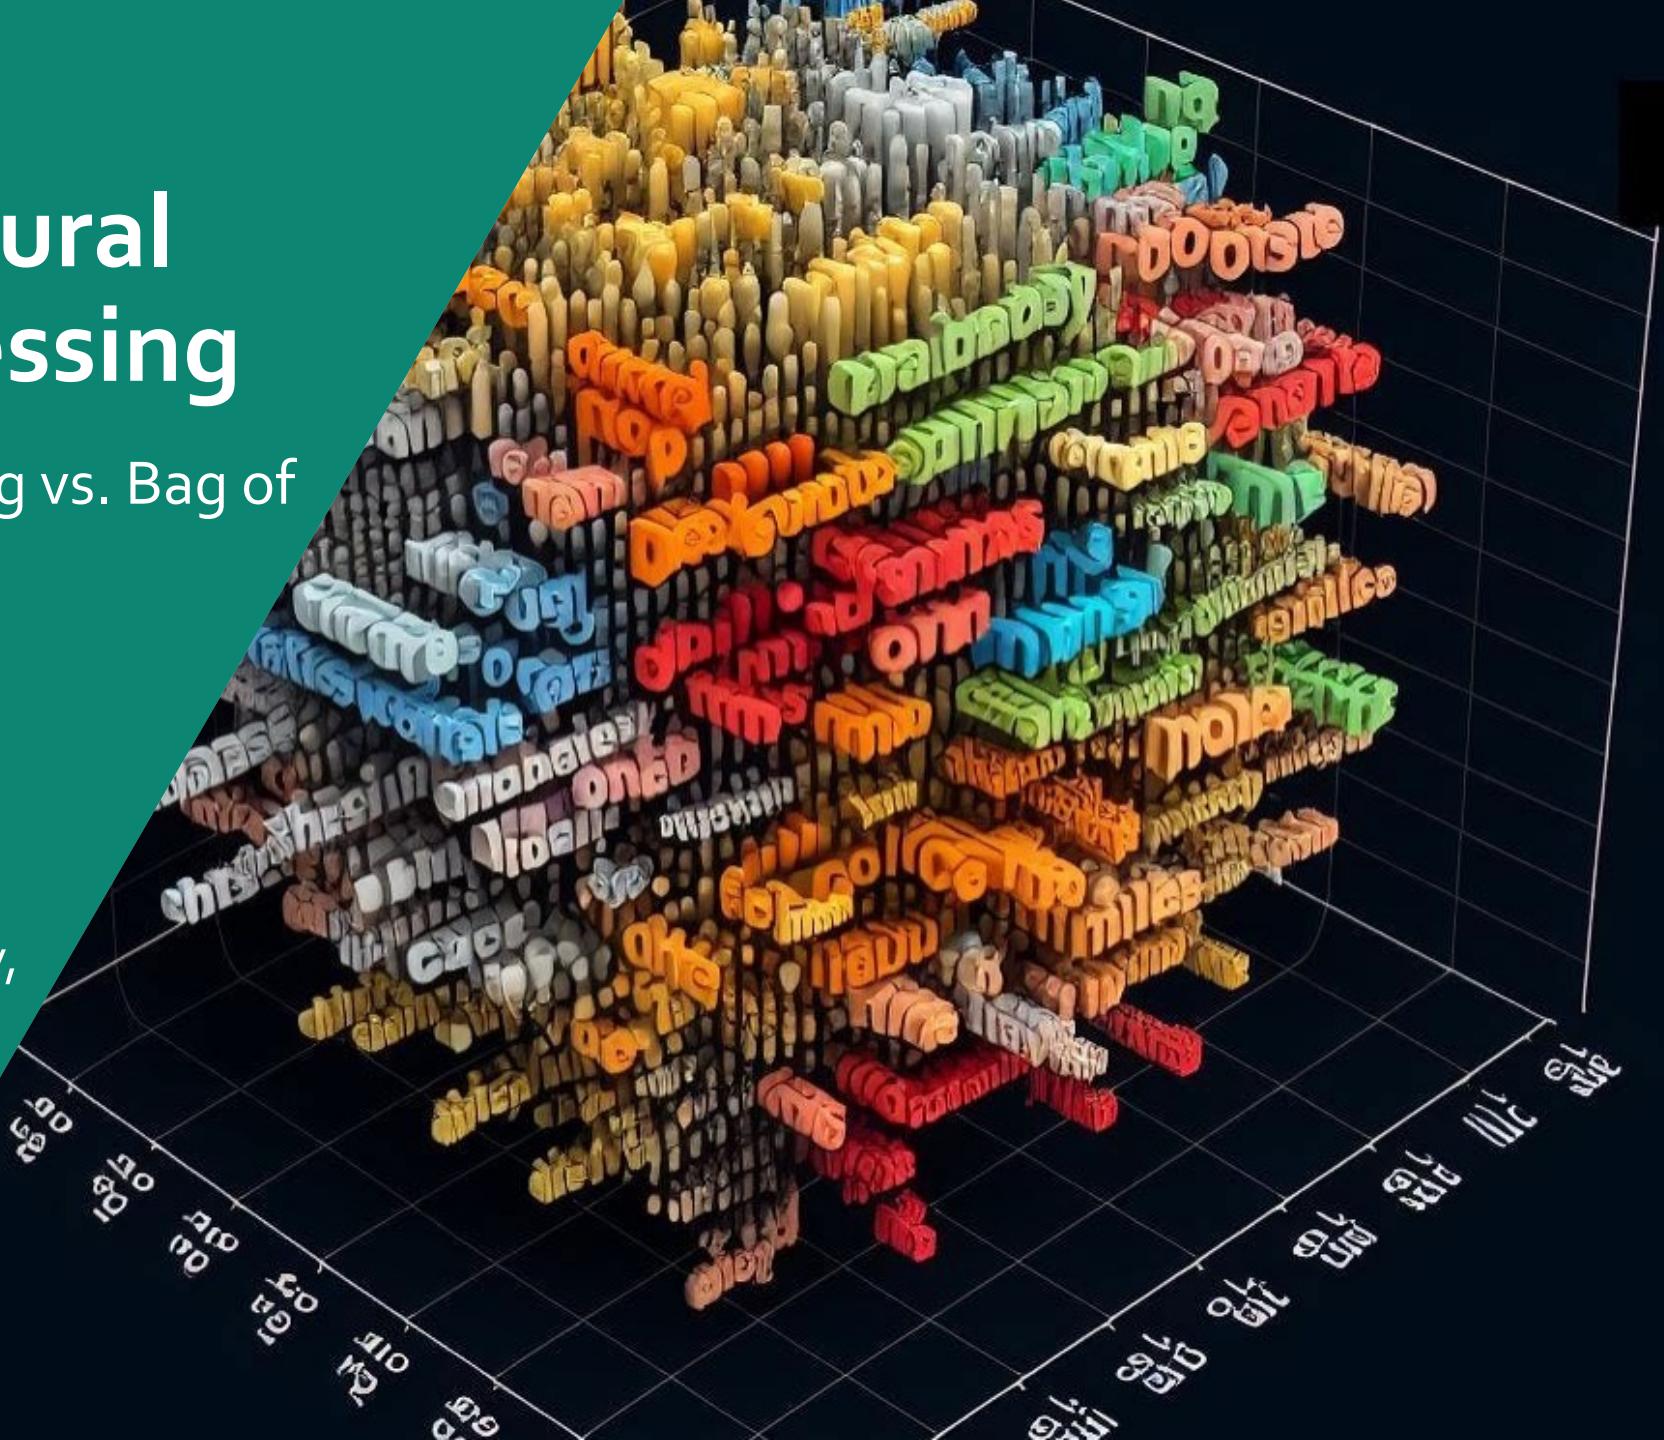
- Start with V random d -dimensional vectors as initial embeddings
- Train a classifier based on embedding similarity
 - Take a corpus and take pairs of words that co-occur as positive examples
 - Take pairs of words that don't co-occur as negative examples
 - Train the classifier to distinguish these by slowly adjusting all the embeddings to improve the classifier performance
 - Throw away the classifier code and keep the embeddings.

SIT330-770: Natural Language Processing

Week 6.6 – Word Embedding vs. Bag of Words

Dr. Mohamed Reda Bouadjenek

School of Information Technology,
Faculty of Sci Eng & Built Env



Traditional Method - Bag of Words Model

Two approaches:

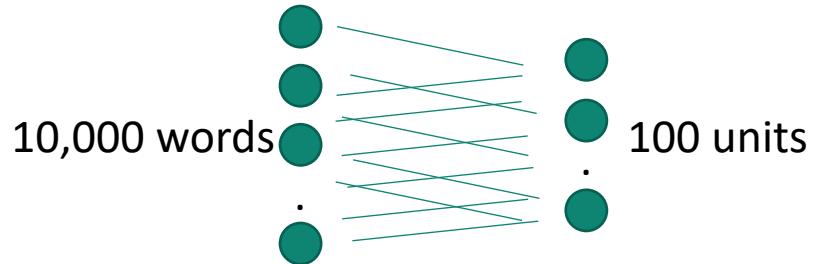
- Either uses one hot encoding.
 - Each word in the vocabulary is represented by one bit position in a HUGE vector.
 - For example, if we have a vocabulary of 10,000 words, and "aardvark" is the *4th word in the dictionary*, it would be represented by: [0 0 0 1 0 0 0 0 0].
- Or uses document representation.
 - Each word in the vocabulary is represented by its presence in documents.
 - For example, if we have a corpus of 1M documents, and "Hello" is in 1th, 3th and 5th documents *only*, it would be represented by: [1 0 1 0 1 0 0 0 0].
- Assumes independence between words.

Word Embeddings

- Stores each word in as a point in space, where it is represented by a dense vector of fixed number of dimensions (generally 300).
 - For example, "Hello" might be represented as : [0.4, -0.11, 0.55, 0.3 . . . 0.1, 0.02].
 - Dimensions are projections along different axes, more of a mathematical concept.
- Unsupervised, built just by reading huge corpus.
- Assumes dependence between words.

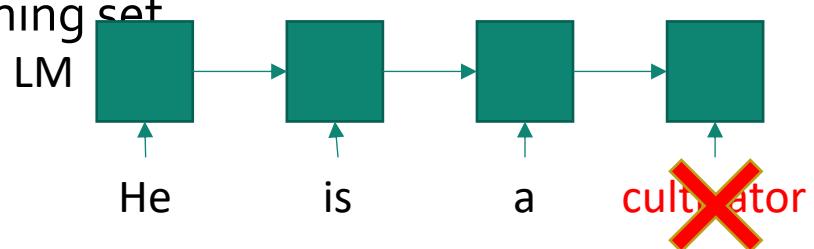
Traditional Method - Bag of Words Model

- Requires **very** large weight matrix for 1st layers.



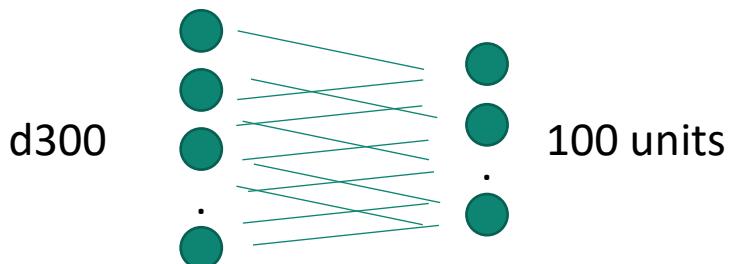
$$W's \ size \ is \ 10,000 \times 100 = 10^6$$

- Models **not flexible** with unseen words in the training set



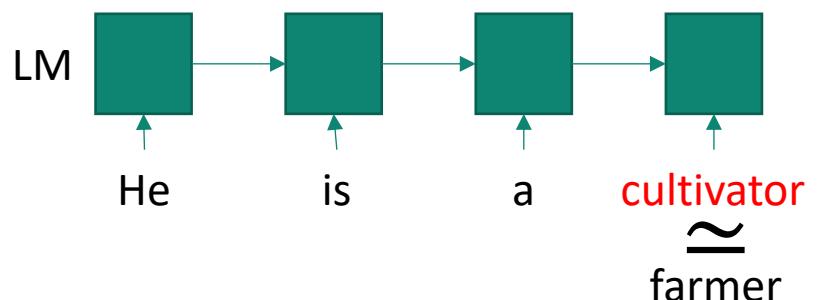
Word Embeddings

- A **compact** weight matrix for 1st layers.



$$W's \ size \ is \ 300 \times 100 = 3 \times 10^4$$

- Flexible** models with unseen words in the training set.

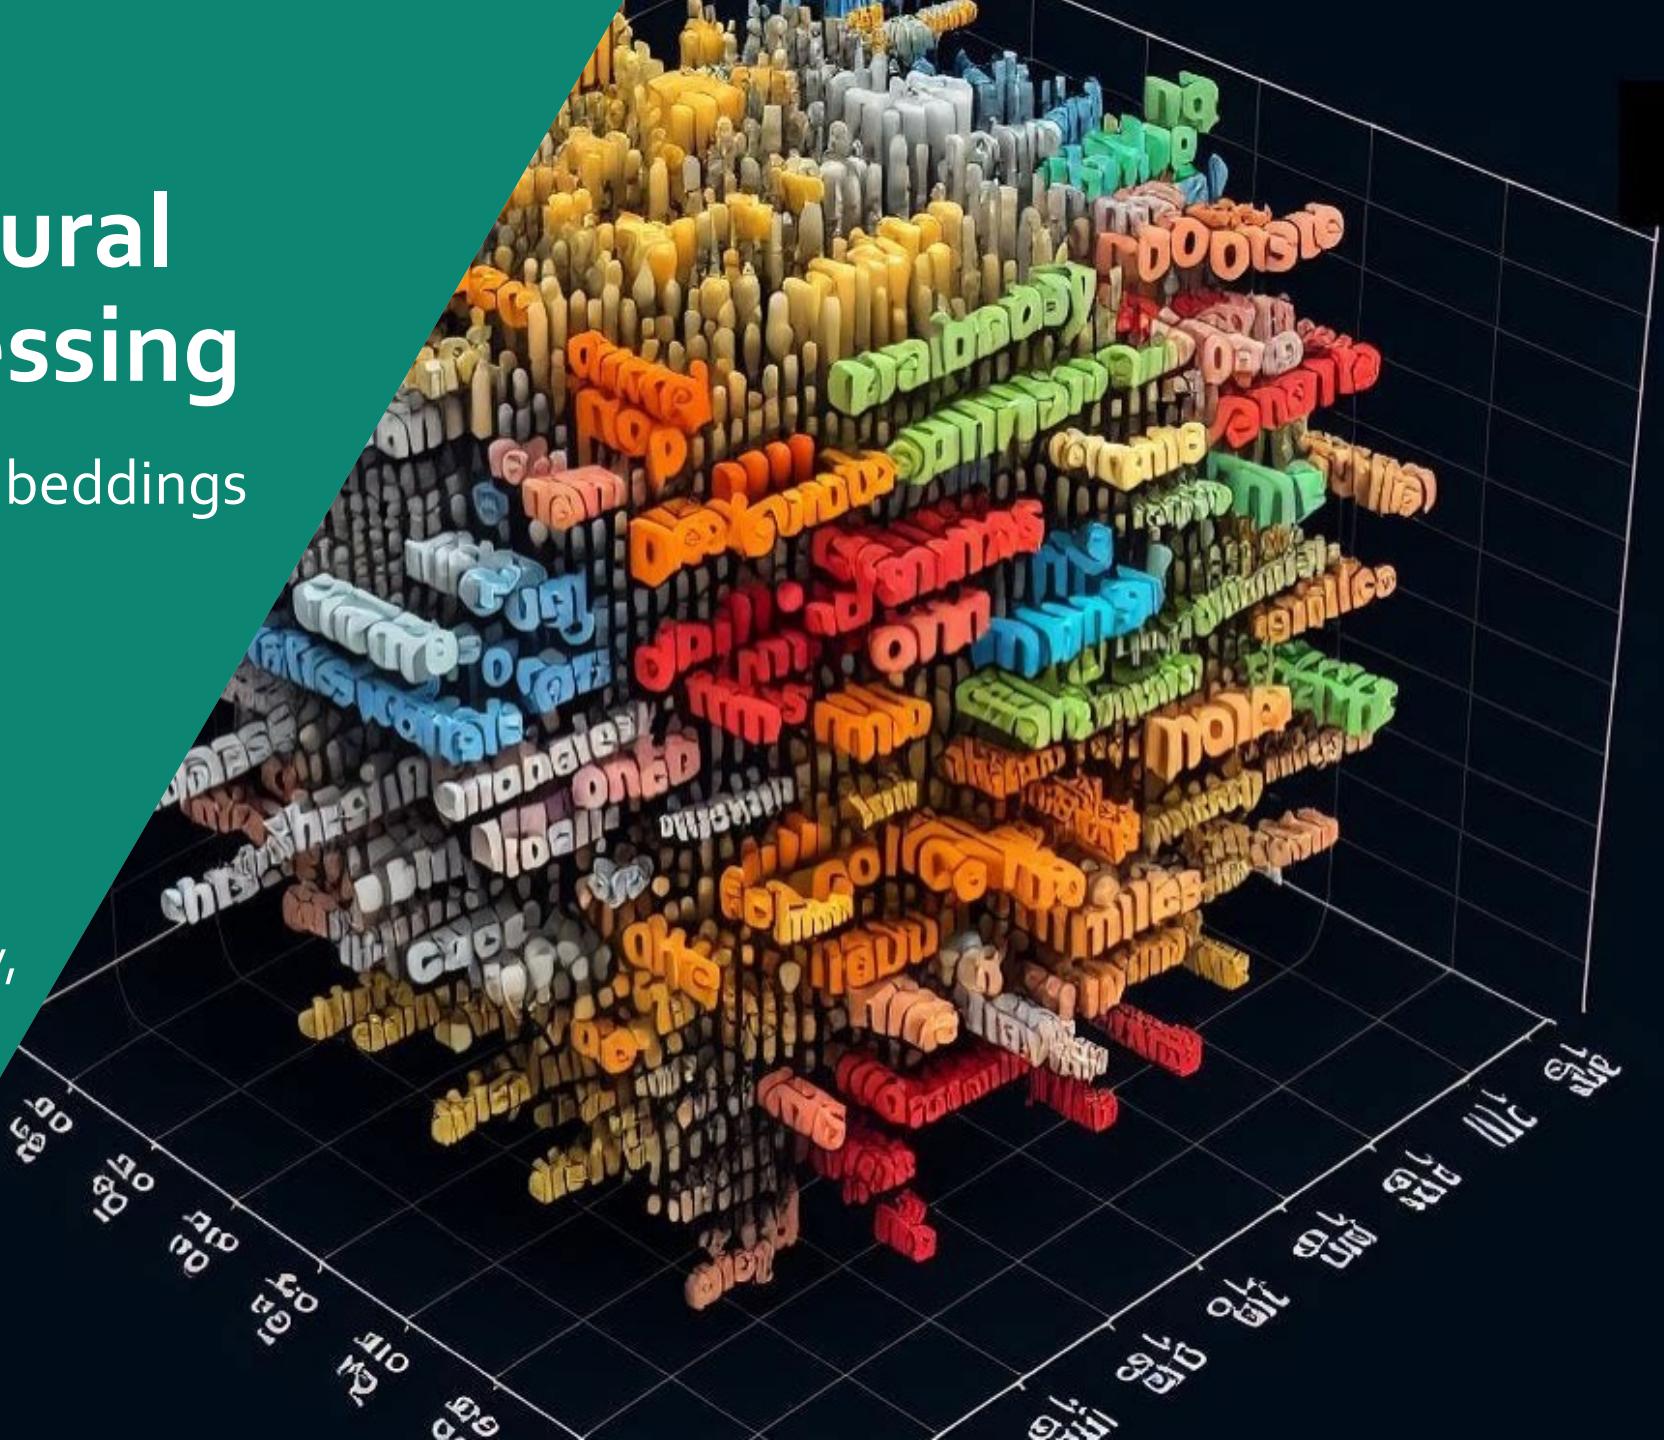


SIT330-770: Natural Language Processing

Week 6.7 – Properties of Embeddings

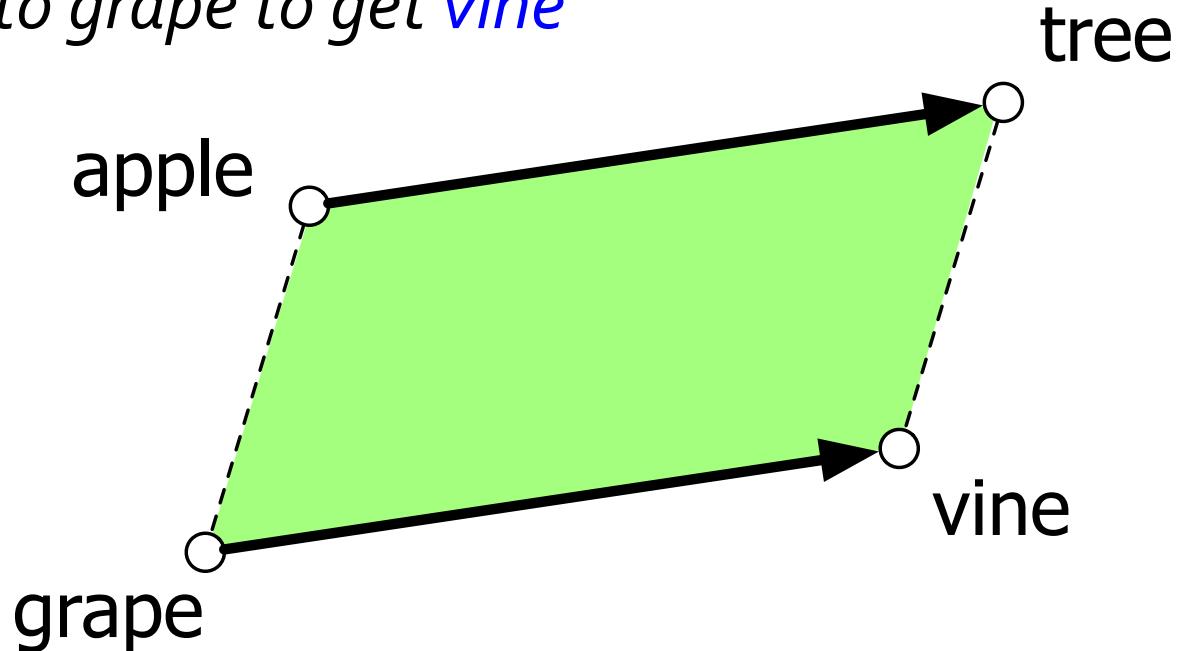
Dr. Mohamed Reda Bouadjenek

School of Information Technology,
Faculty of Sci Eng & Built Env



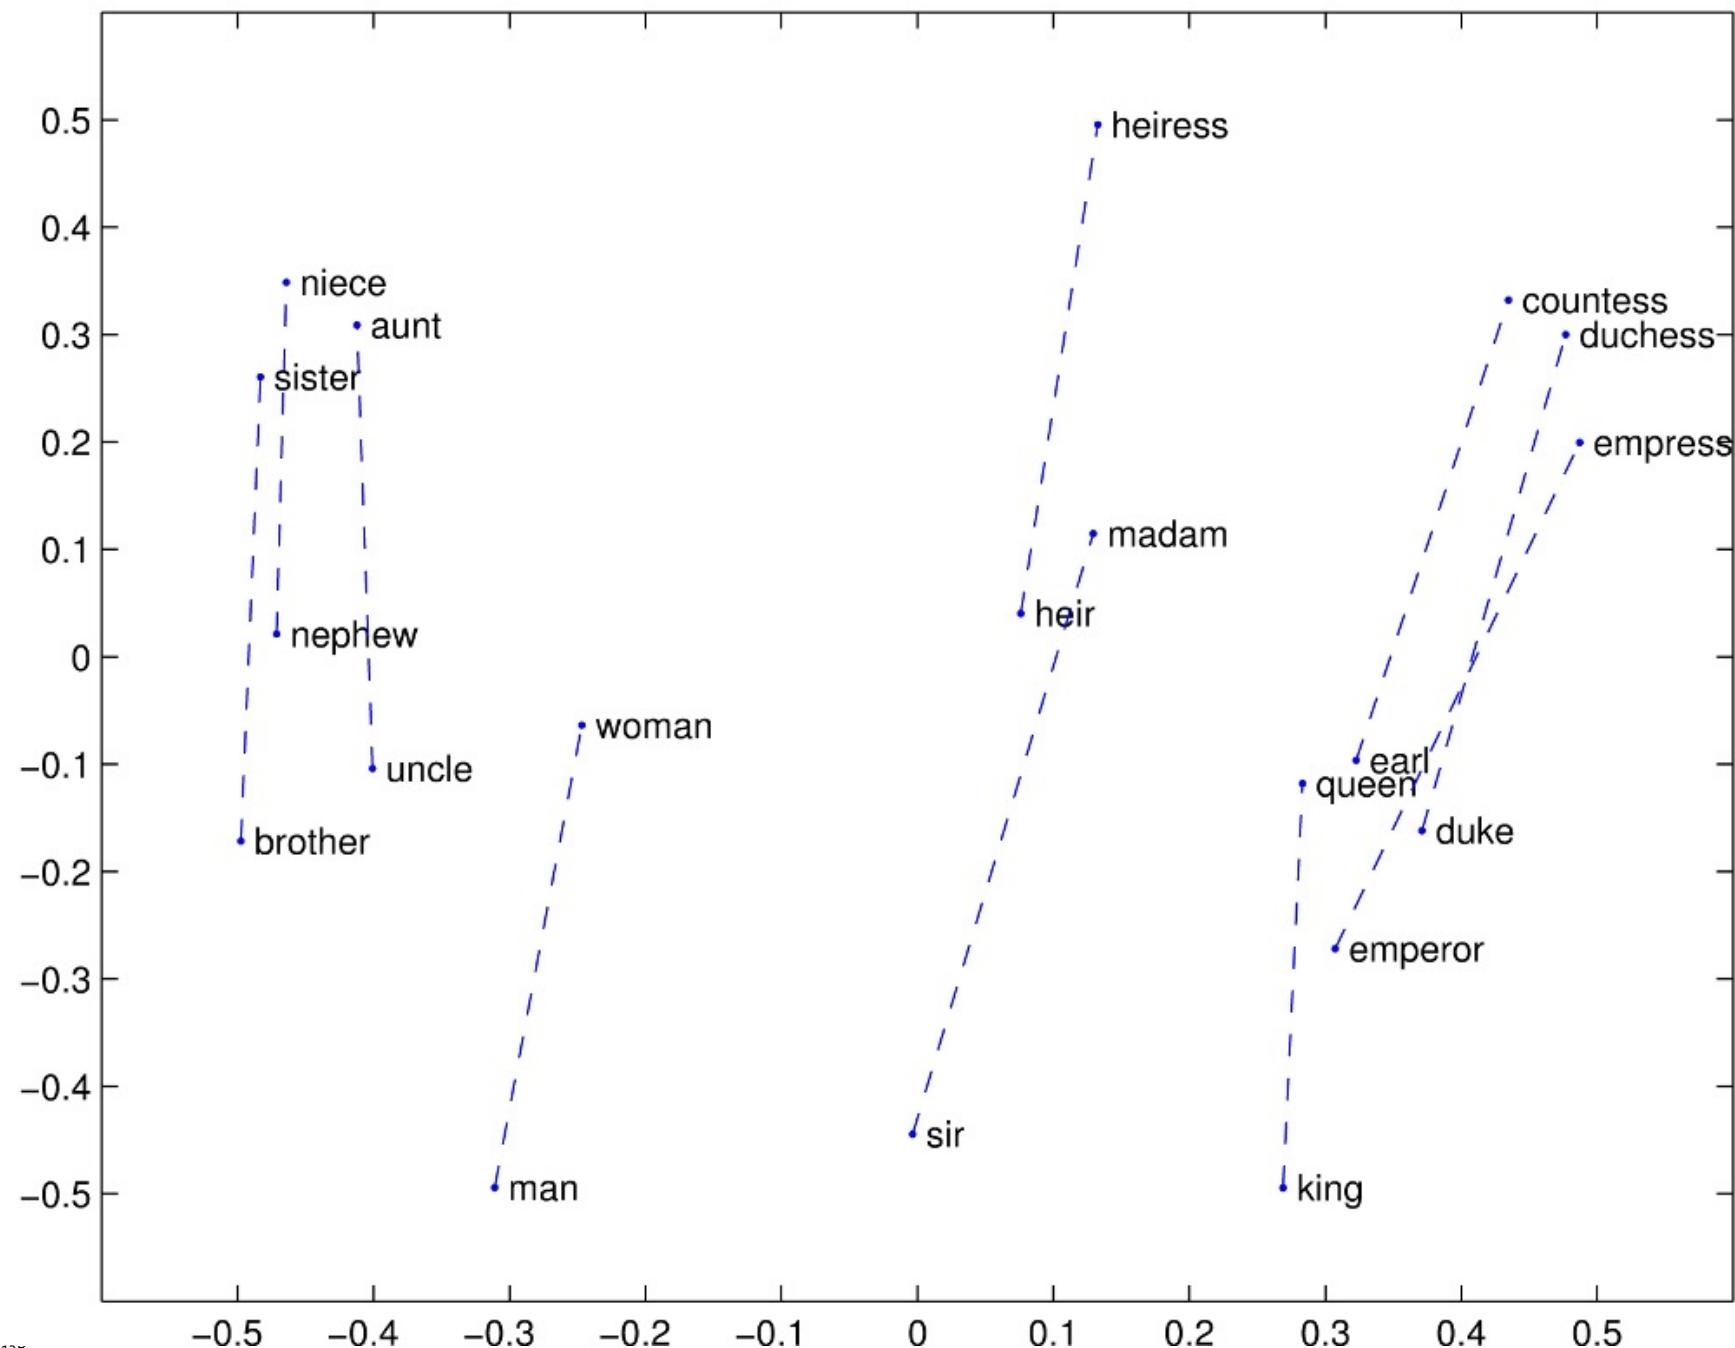
- **Small windows** ($C = +/- 2$) : nearest words are syntactically similar words in same taxonomy
 - *Hogwarts* nearest neighbors are other fictional schools
 - *Sunnydale, Evernight, Blandings*
- **Large windows** ($C = +/- 5$) : nearest words are related words in same semantic field
 - *Hogwarts* nearest neighbors are Harry Potter world:
 - *Dumbledore, half-blood, Malfoy*

- The classic parallelogram model of analogical reasoning (Rumelhart and Abrahamson 1973)
- To solve: "*apple is to tree as grape is to _____*"
- *Add tree – apple to grape to get vine*



- The parallelogram method can solve analogies with both sparse and dense embeddings (Turney and Littman 2005, Mikolov et al. 2013b)
 - king – man + woman is close to queen
 - Paris – France + Italy is close to Rome
- For a problem $a:a^*:b:b^*$, the parallelogram method is:

$$\hat{b}^* = \operatorname*{argmax}_x \text{distance}(x, a^* - a + b)$$



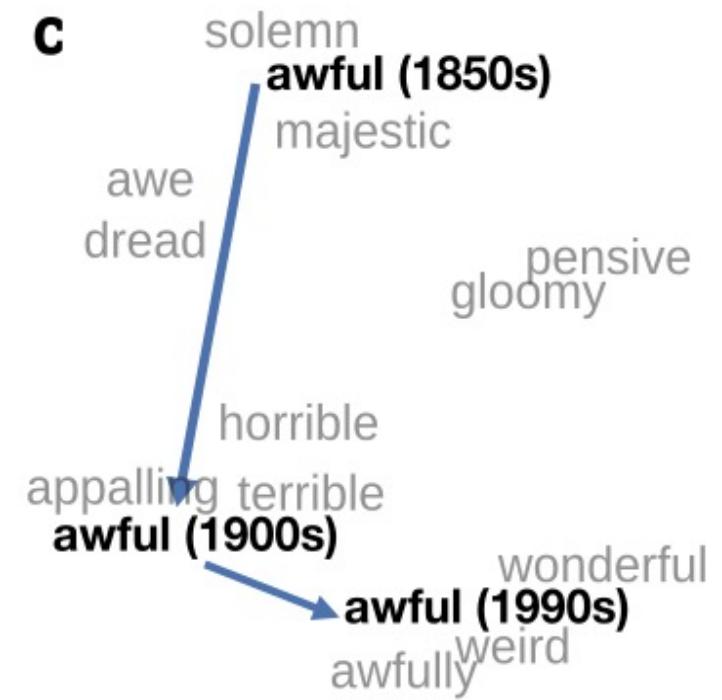
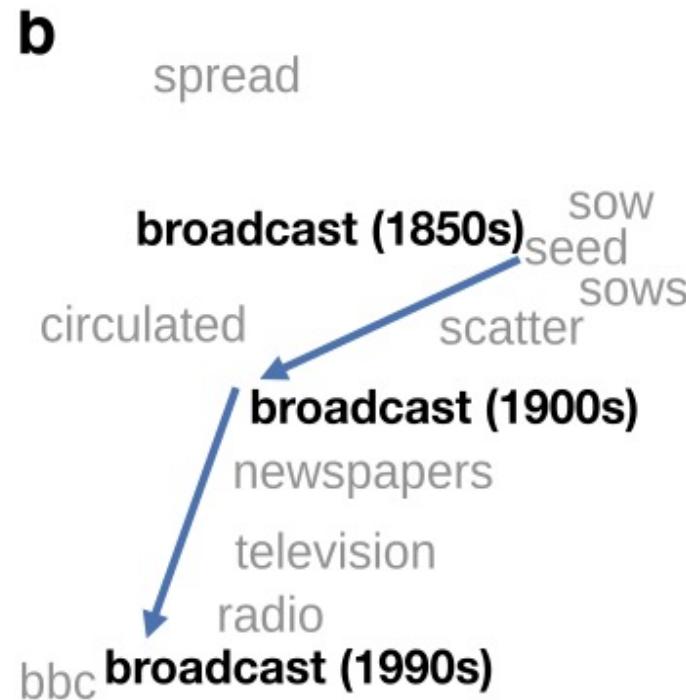
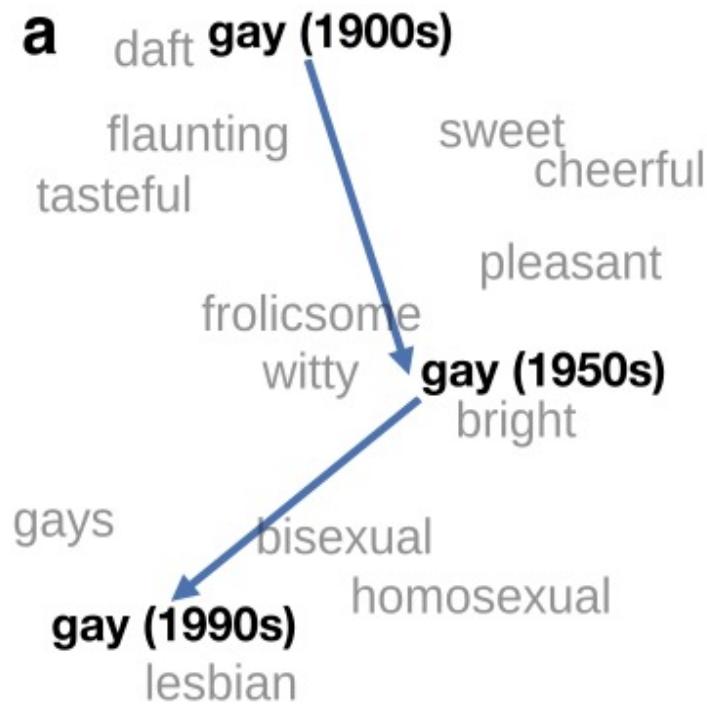
- It only seems to work for frequent words, small distances and certain relations (relating countries to capitals, or parts of speech), but not others.
(Linzen 2016, Gladkova et al. 2016, Ethayarajh et al. 2019a)
- Understanding analogy is an open area of research (Peterson et al. 2020)

Embeddings as a window onto historical semantics



- Train embeddings on different decades of historical text to see meanings shift

~30 million books, 1850-1990, Google Books data



William L. Hamilton, Jure Leskovec, and Dan Jurafsky. 2016. Diachronic Word Embeddings Reveal Statistical Laws of Semantic Change. Proceedings of ACL.

Embeddings reflect cultural bias!



- Ask “Paris : France :: Tokyo : x”
 - x = Japan
- Ask “father : doctor :: mother : x”
 - x = nurse
- Ask “man : computer programmer :: woman : x”
 - x = homemaker

Bolukbasi, Tolga, Kai-Wei Chang, James Y. Zou, Venkatesh Saligrama, and Adam T. Kalai. "Man is to computer programmer as woman is to homemaker? debiasing word embeddings." In *NeurIPS*, pp. 4349-4357. 2016.

Algorithms that use embeddings as part of e.g., hiring searches for programmers, might lead to bias in hiring

Historical embedding as a tool to study cultural biases



Garg, N., Schiebinger, L., Jurafsky, D., and Zou, J. (2018). Word embeddings quantify 100 years of gender and ethnic stereotypes. *Proceedings of the National Academy of Sciences* 115(16), E3635–E3644.

- Compute a **gender or ethnic bias** for each adjective: e.g., how much closer the adjective is to "woman" synonyms than "man" synonyms, or names of particular ethnicities
 - Embeddings for **competence** adjective (*smart, wise, brilliant, resourceful, thoughtful, logical*) are biased toward men, a bias slowly decreasing 1960-1990
 - Embeddings for **dehumanizing** adjectives (barbaric, monstrous, bizarre) were biased toward Asians in the 1930s, bias decreasing over the 20th century.
- These match the results of old surveys done in the 1930s